

Experiment 10

Dimensionality Reduction: Principal Component Analysis (PCA)

Aim:

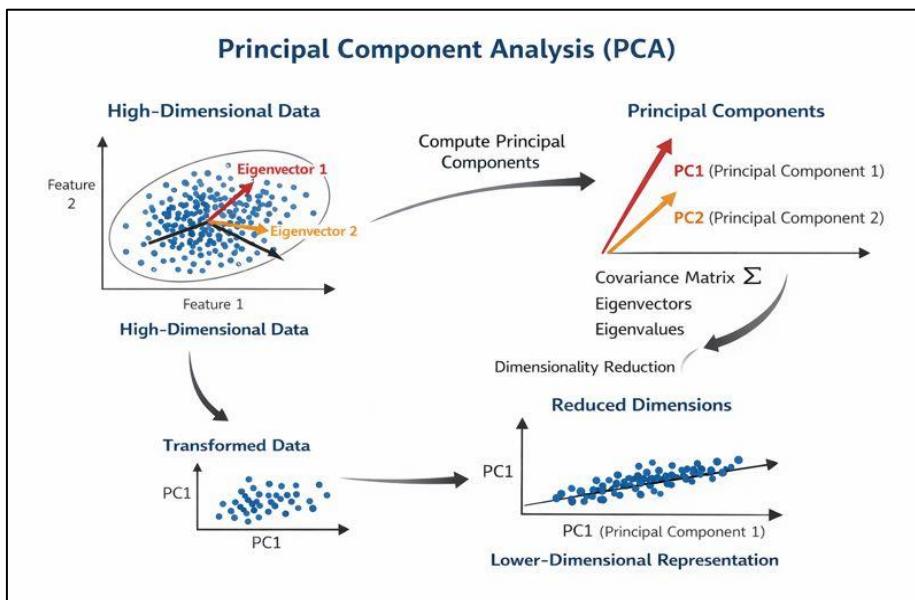
To apply Principal Component Analysis (PCA) for dimensionality reduction by transforming high-dimensional data into a lower-dimensional space while retaining the maximum variance.

Theory:

Principal Component Analysis (PCA) is a dimensionality reduction technique used in unsupervised learning to transform a dataset with many variables into a smaller set that still contains most of the essential information. PCA helps us get rid of non important features. It helps in faster training and inference even data visualization becomes easier. High-dimensional data often includes correlated features, which increase computational complexity and make data interpretation difficult. PCA addresses this problem by transforming the original variables into a new set of uncorrelated variables, called principal components, while retaining as much of the original information as possible.

The objective of PCA is to identify directions in the feature space along which the variance of the data is maximized. The diagram shows how a dataset with many correlated features is first represented in a high-dimensional space. PCA then identifies new axes called principal components that point in the directions of maximum data spread. The original data points are projected onto these new axes, resulting in a reduced-dimensional dataset that still preserves most of the important information.

The mathematical foundation of PCA lies in linear algebra, specifically in eigenvectors and eigenvalues. Eigenvectors define the directions of maximum variance in the data, while eigenvalues indicate the amount of variance explained by each direction.



For a square matrix A , an eigenvector v and its corresponding eigenvalue λ satisfy the equation:

$$Av = \lambda v$$

In PCA, eigenvectors and eigenvalues are derived from the covariance matrix of the data. Eigenvectors associated with larger eigenvalues correspond to more informative principal components. These eigenvectors are normalized to unit length before being used for projection.

By projecting the data onto a smaller number of important principal components, PCA reduces the number of features while keeping most of the useful information. Only the components that capture a large amount of variation in the data are selected. This way, PCA makes the dataset simpler and easier to work with without losing its essential patterns.

Merits of Using PCA:

- Makes large and complex data easier to handle by reducing the number of features
- Removes repeated or similar information from the dataset
- Helps models run faster and more efficiently

Demerits of Using PCA:

- New features created by PCA are hard to understand and explain
- Some useful information may be lost during reduction
- Does not work well when data has non-linear patterns

Code and Result:

Block 1: Importing Required Libraries

INPUT:

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.preprocessing import StandardScaler  
from sklearn.decomposition import PCA  
import ipywidgets as widgets  
from IPython.display import display  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import accuracy_score  
from ipywidgets import interact, IntSlider
```

```
print("Libraries Imported")
```

OUTPUT:

Libraries Imported

Block 2: Reading Data

INPUT:

```
df = pd.read_csv("/content/data.csv")
```

```
print("Data Read Successfully")
```

OUTPUT:

Data Read Successfully

Block 3: Displays the first five rows of the dataset

INPUT:

```
df.head()
```

OUTPUT:

	id	radius_mean	texture_mean	perimeter_mean	..	concave points_worst	symmetry_worst	fractal_dimension_worst	diagnosis
0	842302	17.99	10.38	122.80	..	0.2654	0.4601	0.11890	M
1	842517	20.57	17.77	132.90	..	0.1860	0.2750	0.08902	M
2	84300903	19.69	21.25	130.00	..	0.2430	0.3613	0.08758	M
3	84348301	11.42	20.38	77.58	..	0.2575	0.6638	0.17300	M
4	84358402	20.29	14.34	135.10	..	0.1625	0.2364	0.07678	M

Block 4: Displays the last five rows of the dataset

INPUT:

```
df.tail()
```

OUTPUT:

	id	radius_mean	texture_mean	perimeter_mean	..	concave points_worst	symmetry_worst	fractal_dimension_worst	diagnosis
564	926424	21.56	22.39	142.00	..	0.2216	0.2060	0.07115	M
565	926682	20.13	28.25	131.20	..	0.1628	0.2572	0.06637	M
566	926954	16.60	28.08	108.30	..	0.1418	0.2218	0.07820	M
567	927241	20.60	29.33	140.10	..	0.2650	0.4087	0.12400	M
568	92751	7.76	24.54	47.92	..	0.0000	0.2871	0.07039	B

Block 5:Displays the size of dataset

INPUT:

df.shape

OUTPUT:

(569, 32)

Block 6:Displays summary of dataset, including column names, data types, and non-null counts.

INPUT:

df.info()

OUTPUT:

```
<class 'pandas.core.frame.DataFrame'>

RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):

 #   Column      Non-Null Count Dtype  
 0   id          569 non-null   int64  
 1   radius_mean 569 non-null   float64 
 2   texture_mean 569 non-null   float64 
 3   perimeter_mean 569 non-null   float64 
 4   area_mean    569 non-null   float64 
 5   smoothness_mean 569 non-null   float64 
 6   compactness_mean 569 non-null   float64 
 7   concavity_mean 569 non-null   float64 
 8   concave points_mean 569 non-null   float64 
 9   symmetry_mean 569 non-null   float64 
 10  fractal_dimension_mean 569 non-null   float64 
 11  radius_se     569 non-null   float64 
 12  texture_se    569 non-null   float64 
 13  perimeter_se  569 non-null   float64 
 14  area_se      569 non-null   float64
```

```
15 smoothness_se 569 non-null float64
16 compactness_se 569 non-null float64
17 concavity_se 569 non-null float64
18 concavepoints_se 569 non-null float64
19 symmetry_se 569 non-null float64
20 fractal_dimension_se 569 non-null float64
21 radius_worst 569 non-null float64
22 texture_worst 569 non-null float64
23 perimeter_worst 569 non-null float64
24 area_worst 569 non-null float64
25 smoothness_worst 569 non-null float64
26 compactness_worst 569 non-null float64
27 concavity_worst 569 non-null float64
28 concavepoints_worst 569 non-null float64
29 symmetry_worst 569 non-null float64
30 fractal_dimension_worst 569 non-null float64
31 diagnosis 569 non-null object
dtypes: float64(30), int64(1), object(1)
```

Block 7:Displays the statistics of dataset

INPUT:

```
df.describe()
```

OUTPUT:

Block 8:Counts the frequency of each unique category

INPUT:

```
df['diagnosis'].value_counts()
```

OUTPUT:

Diagnosis	Count
B	357
M	212

Block 9:Separate features and target variable

INPUT:

```
X = df.drop('diagnosis', axis=1)  
y = df['diagnosis']  
print("Feature shape:", X.shape)  
print("Target shape:", y.shape)
```

OUTPUT:

Feature shape: (569, 31)

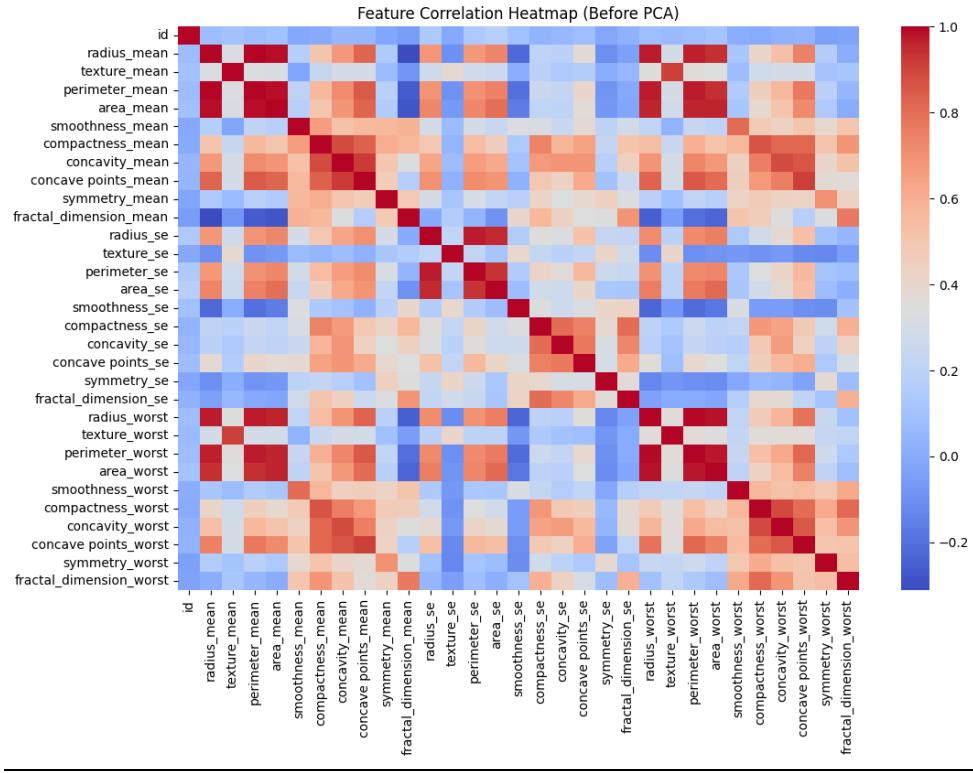
Target shape: (569,)

Block 10:Visualize correlations among features before applying PCA

INPUT:

```
plt.figure(figsize=(12,8))  
sns.heatmap(X.corr(), cmap='coolwarm')  
plt.title("Feature Correlation Heatmap (Before PCA)")  
plt.show()
```

OUTPUT:



Block 11:Standardizing features

INPUT:

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
print("Features Standardized")
```

OUTPUT:

Features Standardized

Block 12:Data Splitting

INPUT:

```
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42, stratify=y
)
print("Data has been Split")
```

OUTPUT:

Data has been Split

Block 13: Apply PCA, and visualize variance

INPUT:

```
pca = PCA(n_components=0.95)

X_train_pca = pca.fit_transform(X_train)

X_test_pca = pca.transform(X_test)

print("PCA applied with 95% variance retained.")
```

OUTPUT:

PCA applied with 95% variance retained.

Block 14: Model training using Logistic Regression

INPUT:

```
clf = LogisticRegression(max_iter=2000)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

acc_before = accuracy_score(y_test, y_pred)

clf_pca = LogisticRegression(max_iter=2000)

clf_pca.fit(X_train_pca, y_train)

y_pred_pca = clf_pca.predict(X_test_pca)

acc_after = accuracy_score(y_test, y_pred_pca)

print("Model Training completed.")
```

OUTPUT:

Model Training completed.

Block 15: Compare classification accuracy before and after applying PCA

INPUT:

```
print("Accuracy BEFORE PCA:", round(acc_before*100, 2), "%")

pca = PCA(n_components=0.95)

X_pca = pca.fit_transform(X_scaled)

print("Reduced Dimensions:", X_pca.shape[1])

X_train_pca, X_test_pca, y_train, y_test = train_test_split(
```

```

X_pca, y, test_size=0.2, random_state=42, stratify=y)

print("Accuracy AFTER PCA:", round(acc_after*100, 2), "%")

plt.figure(figsize=(6,4))

sns.barplot(
    x=['Before PCA', 'After PCA'],
    y=[acc_before, acc_after])

plt.ylabel("Accuracy")

plt.title("Classification Accuracy Comparison")

plt.ylim(0.9, 1.0)

plt.show()

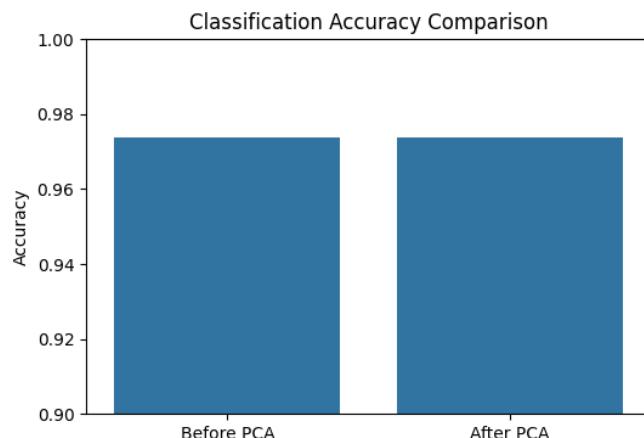
```

OUTPUT:

Accuracy BEFORE PCA: 97.37 %

Reduced Dimensions: 11

Accuracy AFTER PCA: 97.37 %



Block 16:Apply PCA and compute feature loadings for each principal component

INPUT:

```

pca = PCA(n_components=10)

X_pca = pca.fit_transform(X_scaled)

loadings = pd.DataFrame( pca.components_.T, columns=[f"PC{i+1}" for i in range(10)],
index=X.columns)

print("PCA feature loadings calculated for the top 10 principal components.")

```

OUTPUT:

PCA feature loadings calculated for the top 10 principal components.

Block 17: Visualize feature contributions to the first few principal components using a heatmap

INPUT:

```
plt.figure(figsize=(10,8))

sns.heatmap(loadings.iloc[:, :5], cmap="coolwarm",center=0)

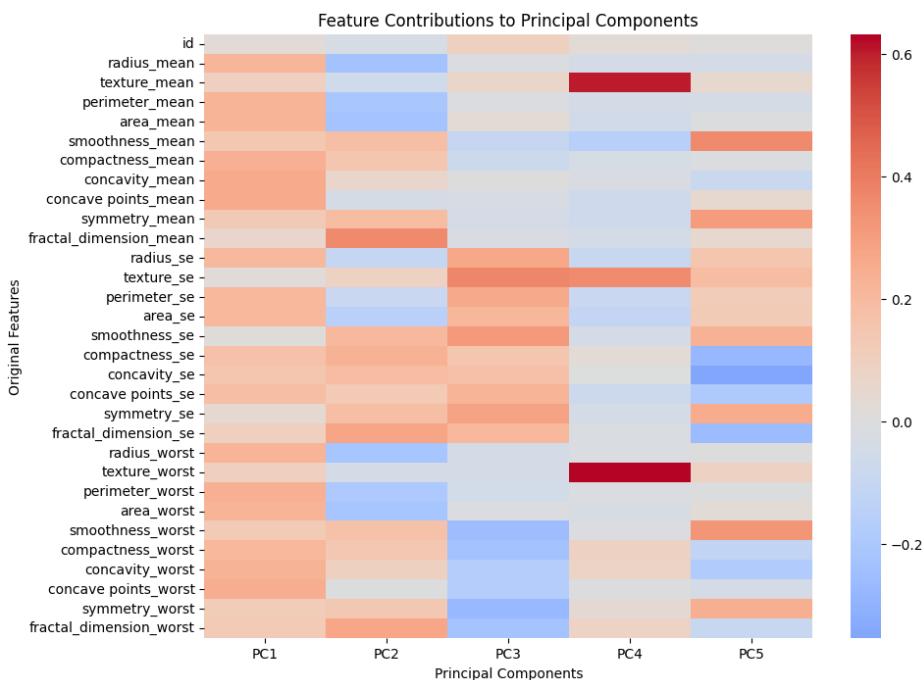
plt.title("Feature Contributions to Principal Components")

plt.xlabel("Principal Components")

plt.ylabel("Original Features")

plt.show()
```

OUTPUT:



Block 18: Interactively analyze and visualize top feature contributions for a selected principal component

INPUT:

```
def interactive_pca(pc_num=1):

    pc = f"PC{pc_num}"

    print(f"Principal Component: {pc}")

    print(f"Explained Variance Ratio: {pca.explained_variance_ratio_[pc_num-1]:.4f}")

    print(f"Cumulative Variance till {pc}: {pca.explained_variance_ratio_[:pc_num].sum():.4f}\n")
```

```

top_features = loadings[pc].abs().sort_values(ascending=False).head(10)

plt.figure(figsize=(8,4))

sns.barplot( x=top_features.values, y=top_features.index)

plt.xlim(0, top_features.max() * 1.1)

plt.xlabel("Feature Contribution Strength")

plt.ylabel("Features")

plt.title(f"Top 10 Feature Contributions to {pc}")

plt.grid(True)

plt.show()

widgets.interact(interactive_pca,
pc_num=widgets.IntSlider(min=1,max=10,step=1,value=1,description="PCA Component"));

```

OUTPUT:

PCA Component= 1

Principal Component: PC1

Explained Variance Ratio: 0.4286

Cumulative Variance till PC1: 0.4286

